

Generalization across model families

critical-token disagreement: best confidence baseline vs TokenTriage · IU X-Ray, 64 reports

■ confidence baseline ■ TokenTriage

Qwen2.5-0.5B

Qwen2.5 · 24 layers · 341 crit. tok


TokenTriage keeps 8% compute saved



SmolLM2-1.7B

SmolLM2 · 24 layers · 349 crit. tok


TokenTriage keeps 83% compute saved



Qwen2.5-1.5B

Qwen2.5 · 28 layers · 341 crit. tok

TokenTriage keeps 33% compute saved



Qwen2.5-3B

Qwen2.5 · 36 layers · 341 crit. tok

TokenTriage keeps 30% compute saved

